

GPU substructure for flashjet

kt / C-A jet substructure on the merge history

exclusive jets · soft-drop grooming · Lund coordinates

Chirayu Gupta

Code: github.com/.../flashjet · branch `benchmarking`, commit `2e912ef`

Reading guide — features: **F1** exclusive-kt · **F2** soft-drop grooming · **F3** Lund ·
inputs: **A** toy events · **B** toy shower · **C** real CMS (C1 QCD, C2 ttbar-2ℓ, C3 ttbar-4q) ·

tree tags `kt` `C/A` `anti-kt`

Summary

Three substructure features, all **pure-torch post-reads of the existing merge history** — no kernel changes, CPU/CUDA identical, negligible cost:

	feature	reads	implements
F1	exclusive-kt jets	kT history	kt algorithm [1111.6097]
F2	soft-drop / mMDT grooming	C/A history	Soft Drop [1402.2657]
F3	Lund coordinates	C/A history	Lund plane [1807.04758]

Validated three ways, each closed against an independent reference:

validation	reference	headline
unit tests	independent NumPy tree-walks	85 passed (13 CUDA-only skipped)
paper closures (toys)	analytic LL predictions	z_g on the $1/z$ curve; areas; β -ordering
real CMS (raw-to-raw)	stored FastJet branches	p_T 1.000000 , m_{SD} -0.004 GeV , R_g 99.2% <0.01

⇒ **flashjet reproduces CMS's FastJet reconstruction to NanoAOD storage precision** — then the same histories become **tagger inputs** (second half of this talk).

What was already there vs. what we added

Pre-existing

- `history.py`: `jet_idx_from_history`, `_resolve_parents`, `_decode`, `splitting_scales_from_history`, `Triton`, `_decode_triton`
- `reference.py`: single-event NumPy ground truth (`cluster_event`)
- `nn_reference.py`: nearest-neighbour reference
- `conftest.py`: `random_event` fixtures
- The kernels + the merge-history recording

The `_ref_*` walkers inside `test_substructure.py` are also ours — independent naive NumPy tree-walks used only to pin the fast implementations.

Added (this work) — in `history.py`

- **F1** `exclusive_jets_from_history`
- **F2** `groom_from_history`
- **F3** `lund_coordinates_from_history`
- helpers: `_resolve_roots`, `_jet_roots`, `_pseudojet_p4`, `_dense_number_roots`
- `api.py`: `ClusterOutput`.{`exclusive_jets`, `groomed_jets`, `mass_drop`, `lund_coordinates`} + a `mask` field
- `tests/test_substructure.py` (new)

The three features

	Feature	Function	Implements
F1	Exclusive jets (kt)	<code>exclusive_jets_from_history(...)</code>	kt algorithm — Fastjet manual [1111.6097], [0802.1189]
F2	Grooming (C/A)	<code>groom_from_history(...)</code>	Soft Drop [1402.2657] · mMDT [1307.0007] · mass-drop [0802.2470]
F3	Lund coordinates	<code>lund_coordinates_from_history(...)</code>	Primary Lund plane [1807.04758]

F1 — undo the last merges of the sequence to expose exactly `n_jets` (or a `d_cut`) subjects; reduces to the inclusive jets at the trivial cut.

F2 — walk each jet down the harder branch, dropping the softer prong until $z > z_{\text{cut}}(\Delta R/R)^\beta$ (soft-drop) — the $O(\log_2 n)$ declustering.

F3 — for every primary split emit $(z, \Delta R, k_t, \ln 1/\Delta R, \ln k_t, d)$; the d -channel ties **exactly** to `splitting_scales` .

How the toys were generated

(the "simulation" — none of it pre-existing; flashjet only clusters, it does not generate events)

Input A: The toy event generators (written in the plot scripts)

flashjet takes 4-vectors in and returns jets — it has **no event generator**. To exercise the features I wrote small, transparent toys *outside the repo*, in the plot scripts.

`make_plots.py` — **two single-fat-jet samples at $p_T \approx 500$ GeV, $R = 0.8$:**

- **QCD-like:** one hard collinear core (92% of p_T , $\sigma_{y,\phi} = 0.06$) + soft wide-angle radiation (8%, $\sigma = 0.30$). Each "spray" = n massive pions with exponential p_T fractions.
- **W-like:** two hard prongs with pair mass $m = \sqrt{z(1-z)} p_T$, $\Delta R = 80.4$ GeV ($z \in [0.30, 0.45]$), random orientation, + 6% soft contamination.

```
def _spray(rng, n, y0, phi0, pt_total, sigma):    # n pions around (y0, phi0)
    frac = rng.exponential(1., n); frac /= frac.sum()
    pt = pt_total * frac
    y, phi = rng.normal(y0, sigma, n), rng.normal(phi0, sigma, n)
    mt = np.sqrt(M_PI**2 + pt**2)
    return np.stack([pt*np.cos(phi), pt*np.sin(phi), mt*np.sinh(y), mt*np.cosh(y)], 1)
```

1500 events per sample. Purpose: QCD vs W is a *known* truth, so any observable that fails to separate them (or lands off the analytic curve) would be a real bug.

Input B: The toy parton shower

`make_paper_plots.py` builds a **primary, fixed-coupling, leading-log** shower — emissions sampled **uniformly in the Lund triangle** with density $\bar{\alpha}$ per unit $(\ln 1/\theta, \ln k_t)$ area, off a hard spine at $(y, \phi) = (0, \pi)$:

$$\theta = e^{-u}, \quad k_t = e^v, \quad z = \frac{k_t}{p_{T0} \theta}, \quad \text{keep } v \leq \ln(p_{T0}/2) - u \text{ } (z \leq \frac{1}{2}), \quad k_t > k_t^{\min}$$

- **Jet areas:** one hard event + a dense grid of infinitely-soft **ghosts** ($p_T = 10^{-8}$); the ghosts each algorithm sweeps up trace its catchment area (the anti-kt-paper construction).

$\bar{\alpha} = 0.25, p_{T0} = 1 \text{ TeV}, R = 0.4, 20\,000$ showers. Everything is re-runnable with a fixed seed (`20260713`).

Input C: the real CMS datasets

Real detector-level **PF-candidate constituents** — the only NanoAOD flavour (JMENano, 150X) that stores the `PFCand` table + the `FatJetPFCand` constituent → AK8 map. Pulled via DAS/ `xrdcp` .

tag	dataset	era	jets	used in
C1	<code>QCD_Pt-15to7000_Flat2018_pythia8</code>	UL18 JMENano (13 TeV)	60 257	CMS (1)–(4), β -family
C2	<code>TTTo2L2Nu</code> (dileptonic)	UL18 JMENano (13 TeV)	12 561	ttbar slides
C3	<code>TTto4Q</code> (fully-hadronic)	RunIII2024 JMENanoV15 (13.6 TeV)	16 516	three-sample compare, tagger inputs

All: leading AK8 jet/event, raw $p_T > 300$ GeV, $|\eta| < 2.4$, 2–200 linked PF candidates.

`PFCand_pt/mass` are **already PUPPI-weighted** (branch titles); stored `FatJet_*` are JEC-corrected

(raw = `pt×(1–rawFactor)`). AK4 uses **MINIAOD** instead (see backup) — no AK4 constituent linker exists in NanoAOD.


Act 1 — Correctness

each feature reproduces the signature figure of the paper it implements

Conclusion of this act: on toy inputs with known truth, F1/F2/F3 land exactly on the papers' analytic predictions.

anti-kt jet shapes — reproduces [0802.1189] Fig. 1

Input B (toy shower), clustering only — no substructure feature




What: each algorithm's **jet catchment area**. kt & C/A give ragged, area-fluctuating jets; **anti-kt** gives rigid **circles** around each hard particle — the defining anti-kt result, from flashjet's own clustering. 1 synthetic event (10 hard particles + ~ 3200 ghosts), $R = 1$.

F1 — kt substructure separates 2-prong from QCD [0802.1189, 1111.6097]

Input A (QCD/W toy fat-jets), kt clustering

kt-based substructure: hard two-prong structure separates from QCD [arXiv:0802.1189, 1111.6097]



Left — $\sqrt{d_{12}}$ (`splitting_scales_from_history`): kt scale of the last merge.
W-like peaks at $\sim m_W/2$, QCD low.

Right — **exclusive 2-subject z**
(`exclusive_jets_from_history`, $n_{\text{jets}} = 2$):
W-like balanced ($z \approx 0.35$), QCD lopsided ($z \rightarrow 0$).

F2 — soft-drop z_g vs the analytic prediction [1402.2657]

Input B (toy shower), C/A soft-drop

The soft-drop momentum fraction z_g from


`groom_from_history`

($z_{\text{cut}} = 0.1$, $\beta = 0$) tracks the **leading-log QCD prediction**

$$p(z_g) = \frac{1/z_g}{\ln(1/2z_{\text{cut}})}$$

across the **entire range**, no free parameters. 72% of jets tagged.

Decisive grooming-correctness plot: the target curve is *the* right answer, so lying on it is unambiguous closure.




F2 — groomed mass ordering in β [1402.2657 Figs. 3–4]

Input B (toy shower), C/A soft-drop at $\beta = 0, 1, 2$

Groomed mass $\rho = m^2/(p_T^2 R^2)$ for $\beta = 0, 1, 2$ vs ungroomed.


- Grooming shifts the spectrum to **lower** ρ (removes soft-wide radiation).
- **Smaller β grooms harder**: the $\beta = 0$ curve reaches furthest left.

Reproduces the β -ordering of the Soft Drop paper.



F3 — primary Lund plane [1807.04758]

Input A (QCD/W toy fat-jets), C/A clustering



`lund_coordinates_from_history` (C/A , $R = 0.8$). **QCD** fills the soft-collinear region smoothly. **W-like** shows the same background **plus an isolated hard-splitting spot** exactly where the 2-body m_W decay must sit (red ★ = predicted position).

Act 2 — On real CMS data (Input C)

not toys any more: real detector-level PF-candidate constituents

Conclusion of this act: on real constituents, flashjet reproduces CMS's stored FastJet output *jet-by-jet, raw-to-raw, to NanoAOD storage precision*.

The real-data pipeline (`make_cms_plots.py`) — dataset C1

These plots use *real jet constituents*, not toys:

1. Pick every CMS **AK8 FatJet** with $p_T > 300$, $|\eta| < 2.4$.
2. Gather its **PF candidates** via
`FatJetPFCand_jetIdx→pfCandIdx→PFCand_{pt, η, φ, m}` .
PFCand_pt/mass are already **PUPPI-weighted** (branch titles) — no separate weight branch.
3. Feed *those constituents* to **our** `cluster(R=0.8, antikt)` — CMS AK8 is anti-kt $R = 0.8$.

4. Leading jet → **F2** soft drop ($z_{\text{cut}} = 0.1, \beta = 0$, declustering the **C/A** tree, as FastJet's SoftDrop does) — and **F3** Lund.
5. Chunked at 3000 jets (torch backend is $O(N^3)$; each AK8 jet is one independent event, so chunking is exact).

The stored CMS values are JEC-corrected: raw jet =

`FatJet_pt × (1 - rawFactor)` ;


`FatJet_msoftdrop` = `m(sub1+sub2)` **with subjet JECs**
(proven from data: $\Delta = +0.0002$ GeV).

CMS (1) — reclustering closes: our p_T vs CMS FatJet_pt — C1

What: feed CMS's own AK8 constituents to **our** anti-kt $R = 0.8$ and compare the reclustered jet p_T to CMS's stored FatJet_pt, **jet-by-jet**.

Result — EXACT: vs the **raw** jet pt
 (FatJet_pt $\times(1-\text{rawFactor})$) the ratio is
median 1.000000, $\sigma = 2.5\times 10^{-4}$ — pure NanoAOD storage precision. The apparent $\sim 6\%$ offset vs the stored value **is the L1L2L3 JEC**, nothing else: CMS stores corrected p_T , we recompute the raw one.

60 257 jets; HTCondor cluster 9087059.




CMS (2) — soft-drop mass matches CMS **EXACTLY** (raw-to-raw) — C1

Raw-to-raw, on the C/A tree (20 065 jets): our F2 soft-drop mass reproduces CMS

`FatJet_msoftdrop` jet-by-jet — the diagonal is exact to NanoAOD storage precision.

	our – CMS
m_{SD}	–0.004 GeV (95.6% <0.5 GeV)
z_g	$ \Delta = 7 \times 10^{-5}$

F2 reproduces CMS `msoftdrop` . The residual 4.4% tail is fully attributed to NanoAOD storage (see backup — soft-candidate table floor + rounding).




CMS (3) — primary Lund plane of 60 257 real QCD jets — C1


What: F3 `lund_coordinates` on the same real AK8 jets.

Needs no comparison curve — it *is* a clean, publication-quality primary Lund plane straight from detector-level simulation.

The full [1807.04758] structure emerges with **no toy input**: the hard-collinear perturbative ridge, the soft plateau, and the three kinematic edges.



CMS (4) — full-event clustering recovers CMS's own AK8 jets — C1



What: the realistic path — feed **every** PF candidate in **the event** to flashjet anti-kt $R = 0.8$

(no `FatJetPFCand` per-jet pre-grouping), then match the resulting inclusive jets to CMS's **stored** AK8 jets by $\Delta R < 0.4$.

Result (7 701 matched jets, $p_T > 300$, raw-to-raw): pt **median 1.0000, 100% within 2%**;
match ΔR **median 0.0019**; mass spectra identical. Full-event flashjet reproduces CMS's own Fastjet AK8 reconstruction to milliradian ΔR — no jets known a priori.

On a **ttbar** sample — dataset C2


a second, independent final state · comparison against the sample's own stored FastJet branches

TTTo2L2Nu UL18 JMENano 150X, 344 k events → 12 561 leading AK8 jets. "Comparison against FastJet"

= against CMS's **stored** FatJet_* / SubJet_* / tau* branches, **not** a FastJet re-run.

TTTo2L2Nu is **dileptonic** → these jets are mostly **b-jets + ISR**, not hadronic top decays.

ttbar — flashjet = stored FastJet branches (raw-to-raw) — C2




pt (our anti-kt $R = 0.8$ vs `FatJet_pt×(1-rawFactor)`):
median **1.000002**, σ 2.2×10^{-4} .

soft-drop mass (our C/A tree vs $m(\text{raw sub}_1 + \text{sub}_2)$):
median Δ **-0.041 GeV**, 94.2% < 0.5 GeV.

The exact match holds on a completely different final state (boosted tops + b-jets), not just QCD —
same conclusion, same NanoAOD-precision residual.

ttbar — our substructure vs stored FastJet observables — C2

ttbar substructure: our F1/F2 observables + the sample's stored FastJet tau ratios



Left (F1): our exclusive-2-subjet $\sqrt{d_{12}}$ on ttbar jets; the m_W/m_t lines are **mass-scale references only** (dileptonic sample has no hadronic decays).

Middle (F2): our soft-drop z_g vs CMS's **raw subjet** z , jet-by-jet — median $|\Delta| = 1.6 \times 10^{-4}$.

Right (context): the sample's stored FastJet **N-subjettiness** ratios — τ_2/τ_1 (W-tag), τ_3/τ_2 (top-tag). Our z_g tracking CMS's stored subjet z to 10^{-4} closes F2 directly against FastJet's SoftDrop output.


Three samples side by side

*QCD (C1) vs dileptonic $t\bar{t}$ (C2) vs **fully-hadronic** $t\bar{t}$ (C3) — the first sample with real boosted W /top decays*

C1 160 393 jets · C2 8 639 jets · C3 (`TTto4Q` RunIII2024 JMENanoV15, 13.6 TeV) 16 516 jets.

All: leading AK8 jet/event, raw $p_T > 300$, $|\eta| < 2.4$, ≤ 200 constituents (`make_compare_plots.py` , HTCondor 9099026).

Lund planes: QCD → b-jets → boosted tops (F3) — C1/C2/C3




The top-decay scale appears exactly where it must: the $t\bar{t} \rightarrow 4q$ plane (C3) grows a hard-splitting

blob at $\ln k_t \approx \ln(m_W/2) \approx 3.7$, wide-angle — over **2×** QCD in the ratio panel (right).

The dileptonic sample (C2, b-jets, no hadronic top) shows only a mild version. Same clustering, same F3, three physics regimes.

Spectra: m_W / m_t peaks from our grooming (F1 + F2) — C1/C2/C3

Sample comparison, leading AK8 jet ($p_T^{raw} > 300$ GeV): QCD vs dileptonic vs fully-hadronic ttbar — all flashjet




Left (F2): our raw soft-drop mass — C3 peaks at m_W with a top shoulder at m_t ; C2 broad and lower;

C1 falls. **Middle (F2):** z_g — ttbar flatter than QCD's $\sim 1/z$. **Right (F1):** $\sqrt{d_{12}} = m\sqrt{z/(1-z)}$ —

W-window jets peak at $39 \approx m_W/2$, top-window at $81 \approx m_t/2$; QCD collapses to low scales.

R_g — a second jet-by-jet exact match (F2 vs stored subjects) — C1/C2/C3



What: our split angle R_g (`groom_from_history` 's `dR`) vs stored $\Delta R(\text{sub}_1, \text{sub}_2)$, **jet-by-jet:** median $\Delta \leq 2 \times 10^{-4}$ — with z_g , both grooming observables close.

Run 3 (C3) = the same mechanism at higher pileup: relative agreement stays **0.14%**.

sample	p_T ratio	$m_{SD} \Delta$	$R_g \text{ } \Delta < 0.01$
C1 QCD UL18	0.999999	-0.004 (95.7%<0.5)	99.2%
C2 $t\bar{t} \text{ } 2\ell 2\nu$	0.999998	-0.033 (94.1%<0.5)	95.9%
C3 $t\bar{t} \text{ } 4q \text{ '24}$	0.999556	-0.077 (69.3%<0.5)	87.0%

Act 3 — From validation to tagger inputs

the same histories that reproduce Fastjet are a free, physics-rich input vector

Conclusion of this act: the 18-var history vector reaches AUC 0.827 for boosted top/W vs QCD — and localizes *where* it helps (boosted decays, not AK4 flavour).

Which tree does each variable come from?

Each jet is clustered **three times** (`extract_tagger_vars.py`); every input is a post-read of one history:

anti-kt $R = 0.8$ — the physical jet

`algorithm="antikt"` → leading-jet 4-vector

- m_{ung} (ungroomed mass), n_{const}

kt **exclusive splitting scales**

`algorithm="kt"` → `splitting_scales()`

- $\sqrt{d_{12}}, \sqrt{d_{23}}, \sqrt{d_{34}}$ (value-sorted)
- $\Rightarrow$ ratios $d_{23}/d_{12}, f_{21}, f_{32}$


C/A **big- R recluster** → **soft-drop + Lund**

`algorithm="cambridge"` → `groomed_jets` + `lund_coordinates`

- groom: $m_{SD}, z_g, R_g, n_{drop}$
- Lund: $n_{Lund}, n(k_t > 1), n(k_t > 5), \ln k_t^{(1,2,3)}, z, \Delta R$ of hardest emission

Why two trees? kt is *value-sorted* — its exclusive scales $\sqrt{d_{12}} \geq \sqrt{d_{23}} \geq \dots$ read off the hardest splittings directly (prong hierarchy). C/A is *angular-ordered* — its primary declustering is the Lund plane / soft-drop sequence. The colored tags below mark each variable's source.


All inputs (1/4) — mass-scale variables



Filled = QCD (pt-reweighted), red = $t\bar{t} \rightarrow 4q$ (C3); per-panel weighted single-variable AUC. All at **0.78–0.79** and 0.8–0.97 correlated — they probe the *same* 2-prong decay mass, so add little to one another.

- **C/A** m_{SD} (0.782): soft-drop mass — clean $m_W \approx 80$ peak + $m_t \approx 160$ shoulder; QCD a steep low-mass continuum.
- **kT** $\sqrt{d_{12}}$ (0.792, best of group): momentum-weighted mass scale of the hardest split.
- **anti-kT** m_{ung} (0.788): ungroomed mass — signal keeps its mass, QCD's is inflated by soft wide radiation.
- **C/A** $k_{t,g} = z_g p_T R_g$ (0.783): k_t of the groomed split, another mass proxy.
- **C/A** $\ln m_{SD}$ (0.782), $\ln \rho$ (0.784): log forms. QCD $\sim$ flat in $\ln \rho$; signal piles at the decay mass.


All inputs (2/4) — prong-hierarchy / kT variables



- **kT** $\sqrt{d_{23}}$ (0.684): 2nd kT scale — populated for 3-prong top ($W \rightarrow q\bar{q}$ inside), near zero for 2-prong W or 1-prong QCD.
- **kT** $\sqrt{d_{34}}$ (0.633): 3rd splitting — weakest, mostly extra radiation.
- **kT** $d_{23}/d_{12} = f_{21}$ (0.655): the ratio removes overall scale — signal peaks near 0.1, QCD broad.
- **kT** $f_{32} = \sqrt{d_{34}/d_{23}}$ (0.628): next ratio in the hierarchy.
- **kT+C/A** $f_z = \sqrt{d_{12}/m_{SD}} = \sqrt{z/(1-z)}$ (0.638): momentum sharing, **mass-decorrelated** — decay shares evenly, QCD soft-biased.

The kT splitting scales *beyond the first* and their **ratios** — do multiple hard prongs exist with a hierarchy (3-prong top, 2-prong W) vs QCD's single DGLAP ladder?


All inputs (3/4) — Lund / counting variables



How many hard emissions, and how hard the sub-leading ones are — the **declustering sequence**, information a single mass number cannot carry. All from the **C/A** primary declustering.

- **C/A** n_{Lund} (0.602), $n(k_t > 1)$ (0.587): primary-emission counts. Signal has *fewer* despite higher pileup — physical.
- **C/A** $n(k_t > 5)$ (0.663): counting only *hard* emissions — top/W give 1–2, QCD more.
- **C/A** $\ln k_t^{(2)}$ (0.688): 2nd-hardest emission — signal bump = the **second decay splitting** (top $\rightarrow W \rightarrow q\bar{q}$).
- **C/A** $\ln k_t^{(3)}$ (0.613): 3rd emission — weaker.
- **C/A** n_{drop} (0.769, **best non-mass var**): soft-drop declustering count. Decays pass in **0–1** steps; QCD needs up to ~12.


All inputs (4/4) — grooming-geometry variables



Geometry of the split soft drop keeps, plus constituent count and grooming survival.

- **C/A** z_g (0.608): groomed momentum share — near the $z_{cut} = 0.1$ edge; modest alone.
- **C/A** R_g (0.779): groomed opening angle, **strong**. Decay angle fixed and wide; QCD collinear.
- **C/A** z (0.698), ΔR (0.714) **of the hardest- k_t emission**: for a decay this *is* the decay split.
- **anti-kT** n_{const} (0.602): constituent multiplicity — quark/gluon-like, weak alone.
- **C/A+anti-kT** $f_{groom} = m_{SD}/m_{ung}$ (0.747): **grooming survival** — decays keep mass (≈ 1), QCD loses it ($\ll 1$).

The result: the declustering *sequence* is the payload





$t\bar{t} \rightarrow 4q$ (C3) vs p_T -reweighted QCD, weighted logistic on the **18-variable** history vector:

- mass-scale vars saturate at AUC **0.794** (they're 0.8–0.97 correlated — the *same* 2-prong mass);
- adding the **declustering sequence** → full set **0.827**, ~2× QCD rejection at 30% eff;
- n_{drop} **alone 0.769** — decay jets pass soft drop in 0–1 steps, QCD needs many;
- $\ln k_t^{(2)}$ resolves the **second** decay splitting (top → W).

The point: this input vector is **free** once the history exists — no new clustering, GPU-batchable. Exploratory (no gen-match, linear model).
Condor 9128460.

What the variables measure: one jet's full merge history

The full merge history of one boosted $t\bar{t} \rightarrow 4q$ jet (25 PF constituents) — red = constituent (leaf), blue = merge node




flashjet stores the **complete** binary tree (`hist_p1, p2, child, d`) — every merge, **nothing removed**. One boosted top jet (25 const., p_T 789, m 135). **Left C/A** with soft-drop overlaid: **green = the groomed jet**, grey = the $n_{\text{drop}} = 5$ dropped soft prongs, ★ = passing split (m_{SD}, z_g, R_g live here). **Right kT** (value-sorted): top merges' $\sqrt{d}R = \sqrt{d_{12}} \geq \sqrt{d_{23}} \geq \dots$. **Grooming = a pruned path through the C/A tree**


...and here grooming works: recovering m_W

Grooming that works: one boosted $t\bar{t} \rightarrow 4q$ jet where soft-drop strips soft junk and recovers $m \approx 83$ GeV (W)

C/A tree — grooming HELPS here
 $m_{\text{ung}} 120 \rightarrow m_{\text{SD}} 83$ GeV (W); $z_g 0.49$, $R_g 0.34$, $n_{\text{drop}} 4$: green spine = groomed jet, grey = dropped



kT tree — FULL history
 (top merges' $\sqrt{d}R = \sqrt{d_{12}} \geq \sqrt{d_{23}} \geq \dots$)



A boosted $W \rightarrow q\bar{q}$ (30 const., p_T 507). Ungroomed mass inflated to **120 GeV** by soft wide radiation; soft-drop peels off 4 soft prongs (grey) and lands the ★ on a **balanced, wide** split — $z_g = 0.49$, $R_g = 0.34$, k_t jumps to **80 GeV** — giving $m_{\text{SD}} = 83 \approx m_W$. Two slides show the same algorithm succeeding vs misfiring — why the tagger uses n_{drop} + the **kT** scales, not m_{SD} alone.

AK8 vs AK4: the history is a boosted-decay tool

Same 18 history variables, best single-variable AUC in three studies:

study	what's being separated	best variable	AUC	reading
AK8 top/W vs QCD	boosted 2-/3-prong decay vs QCD	$\sqrt{d_{12}}$	0.79	history separates strongly
AK4 ttbar vs QCD	single quark jet (dijet) vs QCD	m_{SD}	0.55	one AK4 jet $\approx$ one quark — no in-jet decay
AK4 b vs light	flavour (b vs udsg)	z hardest emis.	0.55	a lifetime question (IP/SV) — invisible to a kinematic tree

⇒ **Pitch history tokens at boosted 2-/3-prong tagging, not as an AK4 / b-tag input** — quantitative confirmation that UParT's impact-parameter / secondary-vertex inputs carry information no clustering tree contains.

Sources: AK8 = MINIAOD `slimmedJetsAK8` (same TTto4Q + QCD Run3 2022 files as AK4); AK4 = MINIAOD `slimmedJets`. Grouped-bar figure `ak4_vs_ak8.png` in backup once same-source AK8 substructure completes.


Summary

flashjet reads the full merge history — F1 exclusive-kt jets, F2 soft-drop grooming, F3 Lund — as pure-torch post-reads, validated to **NanoAOD storage precision** against CMS's FastJet output on QCD and ttbar (raw-to-raw, jet-by-jet).

Those same histories are a free tagger input: the 18-variable vector reaches **AUC 0.827** for boosted top/W vs QCD (the declustering *sequence*, not just mass), and the study cleanly localizes where the history helps — **boosted 2-/3-prong decays**, not AK4 flavour tagging.


Backup

Full-history tree gallery — QCD vs clean top




QCD (23 const, p_T 331): $m_{\text{ung}} 33 \rightarrow m_{\text{SD}} 1.2, n_{\text{drop}} = 13$. Green spine is a long **staircase** — soft prong after soft prong dropped, mass collapses. No balanced hard split. *Why n_{drop} is the best non-mass variable.*

Spine length 13 vs 1 is the entire discriminant, visualised.




Clean top (32 const, p_T 529): $m_{\text{ung}} 156 \rightarrow m_{\text{SD}} 155, n_{\text{drop}} = 1$. Grooming barely works — the ★ sits at the very top on a wide balanced split.

Full-history tree gallery — boosted & b-jet




Boosted / collimated (35 const, p_T 629): $R_g = 0.26$. At high p_T the decay angle shrinks — the hard split moves *down* toward the collinear region, the tree compresses.



b jet (26 const, p_T 403): single hard core, no balanced split — looks QCD-like. Foreshadows the AK4 result: **a kinematic tree does not see flavour**.

AK4 jets — from MINIAOD (constituents NanoAOD doesn't link)




Two-stage pipeline (CMSSW py3.9 and `b_hive` py3.11 are ABI-incompatible):

1. **FWLite** reads `slimmedJets` + PF daughters + `hadronFlavour` → constituent npz
2. `b_hive` runs flashjet C/A + kT + soft-drop + Lund → 18 vars + flavour

The b jet and light jet at the *same* p_T have **near-identical** trees.

JMENano has AK4 `Jet_*` + `hadronFlavour` but **no PF → AK4 linker** (only `FatJetPFCand` for AK8). So AK4 trees are impossible there. Went to **MINIAOD** (`slimmedJets` carry `packedPFCandidates`) via DAS: 12 000 AK4 jets, **3552 b / 6593 udsg**.

AK4 result: no flavour tagging (b vs light)




b vs light (udsg), p_T -reweighted — **every variable 0.50-0.59:**

best AK4 variables	AUC
z of hardest emission	0.591
$\sqrt{d_{12}}$	0.588
m_{SD}	0.573
$n(k_t > 5), f_{32}$	~0.51

b vs light is a **lifetime** question (displaced tracks, SVs) — absent from a kinematic tree. Confirms UParT's IP/SV inputs carry what no tree contains.

AK4 result: ttbar vs QCD (dijet)




ttbar AK4 vs QCD AK4 (MINIAOD, same source) — **every variable 0.50–0.55:**

- best: m_{SD} 0.553, $n(k_t > 1)$ 0.551, $\sqrt{d_{23}}$ 0.549;
- a single AK4 jet in ttbar is mostly **one quark** (a b, or a light quark from $W \rightarrow q\bar{q}$) — no multi-prong decay *inside one jet* for the history to see.

The tiny edge is just the b-jet admixture. Complements the b-vs-light null: the history has no AK4 handle at all.

CMS (2b) — anatomy of the residual m_{SD} tail (the 4.4%)



Every input branch is stored with reduced mantissa (`PFCand_pt` ~ 10 bits $\approx 10^{-3}$,

`SubJet_pt/mass` ~ 9 bits, `SubJet_rawFactor` ~ 5 bits $\approx 2\%$) while CMS ran FastJet at full precision.

Attribution (per-jet tests, 19 695 jets): **50%** — soft candidates **missing from** `FatJetPFCand`

(the table has an effective ~ 0.1 GeV floor; the tail is one-sided, its groomed- p_T deficit correlates with Δm , and $\delta m^2 \approx p_T^{jet} p_T^{lost} \Delta R^2$ matches); **23%** within 3σ of

storage rounding; **20%** rounding-sensitive C/A trees (half-ulp jitter moves them); **7%** genuine

$z \approx z_{cut}$ prong flips (14 $\times$ enriched vs core). **Not fixable from NanoAOD — not an**

algorithm error: relative agreement is $\sim 0.1\%$ everywhere. Scripts: `outliers.py` (HTCondor 9098953).

Soft-drop β -family on real QCD jets (F2) — C1


The toy-shower closure, repeated on **164 292 real CMS**

jets: re-groom the same C/A trees at

$\beta = 0, 1, 2$ ($z_{\text{cut}} = 0.1$) and plot $\rho = m^2/(p_T^2 R^2)$.

- grooming pushes mass **down**; **smaller β grooms harder** — the exact ordering of Soft Drop [1402.2657] Figs. 3–4;
- $\beta = 0$ (mMDT) develops the characteristic flat low- ρ tail.

Grooming re-run 3× on the *same* merge histories (a pure post-read: no re-clustering — the point of the history design).



Correction: AUC tie handling

The AUC routine integrated the ROC **without handling ties**. For discrete counts (>80% of jets share one value) that splits tied jets arbitrarily and **manufactures separation**. Caught because AK4 $n(k_t > 5)$ scored **0.773** while b and light had *identical means in every p_T slice* — impossible.

Fixed by building the ROC on **unique values** (proper Mann–Whitney tie handling). Continuous variables unaffected; corrected AK8 counts:

variable	was	corrected
$n(k_t > 1)$	0.593	0.587
$n(k_t > 5)$	0.691	0.663
n_{drop}	0.765	0.769
f_{match}	0.516	0.648

All mass/geometry variables (m_{SD} 0.782, $\sqrt{d_{12}}$ 0.792, R_g 0.779) **unchanged** — AK8 conclusions stand.

Reproducibility & full-statistics numbers

All CMS plots regenerated **raw-to-raw on the C/A tree** at full statistics
(`make_cms_plots.py` , HTCondor cluster 9087059, 60 257 jets):

observable	comparison	result
jet p_T (QCD)	our anti-kt vs <code>FatJet_pt*(1-rawFactor)</code>	median 1.000000 , σ 2.5×10^{-4}
soft-drop mass (QCD)	our C/A-tree vs $m(\text{raw sub}_1+\text{sub}_2)$	median -0.004 GeV , 95.7% <0.5 GeV
z_g (QCD)	our vs raw subjet z	$ \Delta = 7\times 10^{-5}$
jet p_T (ttbar 2L2Nu)	12 561 leading jets, HTCondor 9098883	median 1.000002 , σ 2.2×10^{-4}
soft-drop mass (ttbar)	our C/A-tree vs $m(\text{raw sub}_1+\text{sub}_2)$	median -0.041 GeV , 94.2% <0.5 GeV
full-event (QCD)	all PFCands $\rightarrow$ anti-kt, ΔR -match	7 701 jets, pt 1.0000 , ΔR med 0.0019
R_g (3 samples)	our groomed <code>dR</code> vs $\Delta R(\text{sub}_1,\text{sub}_2)$	median $\Delta \leq 2\times 10^{-4}$; 99.2/95.9/87.0% <0.01

Residual = NanoAOD float storage throughout. Jets: anti-kt $R = 0.8$; grooming/Lund: big- R **C/A** reclustering (as FastJet).

Function reference (added)

```
exclusive_jets_from_history(hist_p1, hist_p2, hist_child, hist_d, mask,
                           n_jets=None, d_cut=None)           # F1
groom_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4, R,
                  z_cut=0.1, beta=0.0, mu=None)              # F2 (soft-drop / mMDT / mass-drop)
lund_coordinates_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4) # F3 -> (B,J,S,6)

# helpers
_resolve_roots(...)    # pointer-jump roots of an arbitrary sub-forest
_jet_roots(...)        # per-jet root pseudojet id, beam-merge order
_pseudojet_p4(...)     # E-scheme 4-momentum of every pseudojet, keyed by id
_dense_number_roots(...) # compact per-particle jet index
```

API surface: `ClusterOutput.exclusive_jets`, `.groomed_jets`, `.mass_drop`,
`.lund_coordinates`, plus a new `.mask` field to map slots → ids.

Reproducing everything

```
# on lxplus, env b_hive (torch 2.5.1)
cd /eos/home-c/cgupta/flashjet/plots/2026-07-13-substructure
micromamba run -n b_hive python make_plots.py          # F1/F2/F3 toy justification + parity
micromamba run -n b_hive python make_paper_plots.py    # anti-kt/SoftDrop/Lund paper figures
micromamba run -n b_hive python make_cms_plots.py      # real CMS UL18 QCD (C1)
micromamba run -n b_hive python make_ttbar_plots.py    # ttbar (C2)
micromamba run -n b_hive python make_compare_plots.py  # C1/C2/C3 comparisons
# AK4/AK8 from MINIAOD: cmsenv in CMSSW_14_1_0_pre4 + voms-proxy-init, then
#   python3 dump_ak4_constituents.py / dump_ak8_constituents.py (FWLite)
#   micromamba run -n b_hive python ak4_substructure.py / ak4_vs_ak8.py

# tests
cd /eos/home-c/cgupta/flashjet/FlastJetDemo
micromamba run -n b_hive python -m pytest -q          # 85 passed, 13 skipped
```

Fixed seed `20260713` throughout. Papers catalogued in `References/Flashjet/papers.md`.