

GPU substructure for flashjet

kt / C-A jet substructure on the merge history

exclusive jets · soft-drop grooming · Lund coordinates

Chirayu Gupta — for Alexandre De Moor

2026-07-13 · branch `benchmarking` (pushed: commit `2e912ef`)

Every claim below is closed against an independent reference and/or the defining paper. All plots regenerable: `plots/2026-07-13-substructure/` .

Where this sits in flashjet

flashjet clusters padded `(B, N, 4)` torch tensors into jets on the GPU, staying on-device for ML training loops. Generalized-kt:

$$d_{ij} = \min(k_{t,i}^{2p}, k_{t,j}^{2p}) \frac{\Delta R_{ij}^2}{R^2}, \quad p = -1 \text{ (anti-} k_t), 0 \text{ (C/A)}, +1 \text{ (} k_t)$$

The key object Alex's kernels already produce: the *merge history*. Every recombination step records four arrays — `hist_p1`, `hist_p2` (the two inputs), `hist_child` (the output id), `hist_d` (the distance at which they merged). That history *is* the full binary clustering tree.

Our contribution = read that tree. No new kernels, no changes to the clustering path. All three features are pure-torch post-reads of `(hist_p1, hist_p2, hist_child, hist_d)`, so they run unchanged on CPU or CUDA and add negligible cost.

What was already there vs. what we added

Pre-existing (Alex)

- `history.py`: `jet_idx_from_history`, `_resolve_parents`, `_decode`, `splitting_scales_from_history`, `Triton`, `_decode_triton`
- `reference.py`: single-event NumPy ground truth (`cluster_event`)
- `nn_reference.py`: nearest-neighbour reference
- `conftest.py`: `random_event` fixtures
- The kernels + the merge-history recording

The `_ref_*` walkers inside `test_substructure.py` are also ours — independent naive NumPy tree-walks used only to pin the fast implementations.

Added (this work) — in `history.py`

- **F1** `exclusive_jets_from_history`
- **F2** `groom_from_history`
- **F3** `lund_coordinates_from_history`
- helpers: `_resolve_roots`, `_jet_roots`, `_pseudojet_p4`, `_dense_number_roots`
- `api.py`: `ClusterOutput.{exclusive_jets, groomed_jets, mass_drop, lund_coordinates}` + a `mask` field
- `tests/test_substructure.py` (new)

The three features

	Feature	Function	Implements
F1	Exclusive jets (kt)	<code>exclusive_jets_from_history(...)</code>	kt algorithm — Fastjet manual [1111.6097], [0802.1189]
F2	Grooming (C/A)	<code>groom_from_history(...)</code>	Soft Drop [1402.2657] · mMDT [1307.0007] · mass-drop [0802.2470]
F3	Lund coordinates	<code>lund_coordinates_from_history(...)</code>	Primary Lund plane [1807.04758]

F1 — undo the last merges of the sequence to expose exactly `n_jets` (or a `d_cut`) subjects; reduces to the inclusive jets at the trivial cut.

F2 — walk each jet down the harder branch, dropping the softer prong until $z > z_{\text{cut}}(\Delta R/R)^\beta$ (soft-drop) — the $O(\log_2 n)$ declustering.

F3 — for every primary split emit $(z, \Delta R, k_t, \ln 1/\Delta R, \ln k_t, d)$; the d -channel ties **exactly** to `splitting_scales` .

How the toys were generated

(the "simulation" — none of it pre-existing; flashjet only clusters, it does not generate events)

The toy event generators (written in the plot scripts)

flashjet takes 4-vectors in and returns jets — it has **no event generator**. To exercise the features I wrote small, transparent toys *outside the repo*, in the plot scripts.

`make_plots.py` — **two single-fat-jet samples at $p_T \approx 500$ GeV, $R = 0.8$:**

- **QCD-like:** one hard collinear core (92% of p_T , $\sigma_{y,\phi} = 0.06$) + soft wide-angle radiation (8%, $\sigma = 0.30$). Each "spray" = n massive pions with exponential p_T fractions.
- **W-like:** two hard prongs with pair mass $m = \sqrt{z(1-z)} p_T$, $\Delta R = 80.4$ GeV ($z \in [0.30, 0.45]$), random orientation, + 6% soft contamination.

```
def _spray(rng, n, y0, phi0, pt_total, sigma):    # n pions around (y0,phi0)
    frac = rng.exponential(1., n); frac /= frac.sum()
    pt = pt_total * frac
    y, phi = rng.normal(y0, sigma, n), rng.normal(phi0, sigma, n)
    mt = np.sqrt(M_PI**2 + pt**2)
    return np.stack([pt*np.cos(phi), pt*np.sin(phi), mt*np.sinh(y), mt*np.cosh(y)], 1)
```

1500 events per sample. Purpose: QCD vs W is a *known* truth, so any observable that fails to separate them (or lands off the analytic curve) would be a real bug.

The toy parton shower (for the paper-figure closures)

`make_paper_plots.py` builds a **primary, fixed-coupling, leading-log** shower — emissions sampled **uniformly in the Lund triangle** with density $\bar{\alpha}$ per unit $(\ln 1/\theta, \ln k_t)$ area, off a hard spine at $(y, \phi) = (0, \pi)$:

$$\theta = e^{-u}, \quad k_t = e^v, \quad z = \frac{k_t}{p_{T0} \theta}, \quad \text{keep } v \leq \ln(p_{T0}/2) - u \quad (z \leq \tfrac{1}{2}), \quad k_t > k_t^{\min}$$

This is **exactly the semi-classical picture the papers' analytic predictions are derived in** — which is what makes the closure *quantitative*, not just qualitative.

- **Jet areas**: one hard event + a dense grid of infinitely-soft **ghosts** ($p_T = 10^{-8}$);
the ghosts each algorithm sweeps up trace its catchment area (the anti-kt-paper construction).

$\bar{\alpha} = 0.25, p_{T0} = 1 \text{ TeV}, R = 0.4$, 20 000 showers. Everything is re-runnable with a fixed seed (`20260713`).


Correctness

each feature reproduces the signature figure of the paper it implements

Every slide states its **input**. Three input classes appear:

- (A) ad-hoc QCD/W toys (`make_plots.py`),
- (B) a leading-log toy shower (`make_paper_plots.py`),
- (C) real CMS PF-candidate constituents (`make_cms_plots.py`). All seed `20260713` .

anti-kt jet shapes — reproduces [0802.1189] Fig. 1




What: each algorithm's **jet catchment area**. kt & C/A give ragged, area-fluctuating jets; **anti-kt** gives rigid **circles** around each hard particle — the defining anti-kt result, from flashjet's own clustering.

How made — input (B), no substructure feature, clustering only.

1 synthetic event: **10 hard massless particles** (p_T 10–400 GeV, random y, ϕ) + a **uniform grid of ~3200 "ghosts"** ($p_T = 10^{-8}$, spacing 0.125 in y, ϕ). Clustered at $R = 1$ with `cluster_event_nn` for $p = +1/0/ - 1$; each ghost is coloured by the jet it lands in.

F1 — kt substructure separates 2-prong from QCD [0802.1189, 1111.6097]



Left — $\sqrt{d_{12}}$ (`splitting_scales_from_history`): kt scale of the last merge.

W-like peaks at $\sim m_W/2$, QCD low.

Right — **exclusive 2-subjet** z

(`exclusive_jets_from_history` , $n_{jets} = 2$):

W-like balanced ($z \approx 0.35$), QCD lopsided ($z \rightarrow 0$).

How made — input (A), leading jet, kt clustering.

1500 **QCD-like** + 1500 **W-like** toy fat-jets ($p_T \approx 500$ GeV, $R = 0.8$); each a bag of

massive-pion 4-vectors from `_spray` — **synthetic particles, not real**

constituents. W pair-mass

= 80.4 GeV. Both observables read the **kt** merge history.

F2 — soft-drop z_g vs the analytic prediction [1402.2657]

The soft-drop momentum fraction z_g from

`groom_from_history`


($z_{\text{cut}} = 0.1$, $\beta = 0$) tracks the **leading-log QCD prediction**

$$p(z_g) = \frac{1/z_g}{\ln(1/2z_{\text{cut}})}$$

across the **entire range**, no free parameters. 72% of jets tagged.

Decisive grooming-correctness plot: the target curve is *the* right answer, so lying on it is unambiguous closure.

How made — input (B). 20 000 toy-shower jets ($p_{T0} = 1$ TeV, $R = 0.4$), leading jet per event, soft-dropped by F2; the black curve is the analytic $1/z$ with *no* fit.



F2 — groomed mass ordering in β [1402.2657 Figs. 3-4]

Groomed mass $\rho = m^2/(p_T^2 R^2)$ for $\beta = 0, 1, 2$ vs ungroomed.


- Grooming shifts the spectrum to **lower** ρ (removes soft-wide radiation).
- **Smaller β grooms harder**: the $\beta = 0$ curve reaches furthest left.

Reproduces the β -ordering of the Soft Drop paper.


How made — input (B). Same 20 000 toy-shower jets; F2

`groom_from_history` re-run at

$\beta = 0, 1, 2$ (all $z_{\text{cut}} = 0.1$) on the leading jet, ρ from the groomed 4-vector.



F3 — primary Lund plane [1807.04758]



`lund_coordinates_from_history` (C/A , $R = 0.8$). **QCD** fills the soft-collinear region smoothly. **W-like** shows the same background **plus an isolated hard-splitting spot** exactly where the 2-body m_W decay must sit (red ★ = predicted position).

How made — input (A). Same 1500+1500 QCD/W toy fat-jets as the F1 slide, **C/A** clustered; F3 emits $(\ln 1/\Delta R, \ln k_t)$ for every primary split of the leading jet, histogrammed over all jets.

On real CMS data — input (C)

not toys any more: real detector-level PF-candidate constituents

The real-data pipeline (`make_cms_plots.py`)

Sample: `QCD_Pt-15to7000_Flat2018_pythia8` , **UL18 JMENano** (150X reprocessing) — the one NanoAOD flavour that stores the `PFCand` table + the `FatJetPFCand` constituent → AK8 map. Pulled via `DAS/ xrdcp` → `data/qcd_jmenano_150x.root` (1.7 GB). Ordinary NanoAOD has no constituents and is useless here.

These plots use *real jet constituents*, not toys:

1. Pick every CMS **AK8 FatJet** with $p_T > 300$, $|\eta| < 2.4$.
2. Gather its **PF candidates** via `FatJetPFCand_jetIdx→pfCandIdx→PFCand_{pt,η,φ,m}` .
`PFCand_pt/mass` **are already PUPPI-weighted** (the branch titles say so) — the weighted 4-vectors are stored directly, which is why no separate weight branch exists.
3. Feed *those constituents* to **our** `cluster(R=0.8, antikt)` — CMS AK8 is anti-kt $R = 0.8$.

4. Leading jet → **F2** soft drop ($z_{\text{cut}} = 0.1, \beta = 0$, declustering the **C/A** tree, as FastJet's SoftDrop does) — and **F3** Lund.
5. Chunked at 3000 jets (torch backend is $O(N^3)$; each AK8 jet is one independent event, so chunking is exact).

The stored CMS values are JEC-corrected: raw jet =

`FatJet_pt × (1 - rawFactor)` ;

`FatJet_msoftdrop` = $m(\text{sub1} + \text{sub2})$ **with subjet JECs** (proven from data: $\Delta = +0.0002$ GeV).


CMS (1) — reclustering closes: our p_T vs CMS FatJet_pt

What: feed CMS's own AK8 constituents to **our** anti-kt $R = 0.8$ and compare the reclustered jet p_T to CMS's stored FatJet_pt, **jet-by-jet**.


Result — EXACT: vs the **raw** jet pt (FatJet_pt $\times(1-\text{rawFactor})$) the ratio is **median 0.999999**, $\sigma = 2\times 10^{-4}$ — pure NanoAOD storage precision. The apparent $\sim 6\%$ offset vs the stored value **is the L1L2L3 JEC**, nothing else: CMS stores corrected p_T , we recompute the raw one.

Input (C). 2D p_T - p_T histogram + mass spectra, real constituents; exactness numbers

from `exact_match2.py` (20 079 jets).



CMS (2) — soft-drop mass matches CMS EXACTLY (raw-to-raw)



The "offset" was never PUPPI — the constituents are *already* PUPPI-weighted. Two stacked artefacts, both removed: (1) JEC — `FatJet_pt / msoftdrop` are JEC-corrected; raw jet = `pt*(1-rowFactor)`, `msoftdrop` = $m(\text{JEC-corrected sub}_1 + \text{sub}_2)$, proven from data ($\Delta = +0.0002$ GeV).

(2) Wrong tree — we groomed the anti-kt history; Fastjet SoftDrop declusters a C/A reclustering.

Raw-to-raw, C/A tree (20 065 jets):

	our – CMS
p_T ratio	0.999999 ($\sigma 2 \times 10^{-4}$)
m_{SD}	-0.004 GeV (95.6% <0.5 GeV)
z_g	$ \Delta = 7 \times 10^{-5}$

Residual = NanoAOD float precision. F2 reproduces CMS `msoftdrop`.


CMS (3) — primary Lund plane of 60 285 real QCD jets

What: F3 `lund_coordinates` on the same real AK8 jets.

Needs no comparison curve — it *is* a clean, publication-quality primary Lund plane straight from detector-level simulation.

The full [1807.04758] structure emerges with **no toy input**: the hard-collinear perturbative ridge, the soft plateau, and the three kinematic edges.

Input (C). 60 285 AK8 jets, $p_T > 300$ GeV; every primary split of each leading jet histogrammed.



Backup

Open follow-ups (CMS plots)

The **CMS (2)** exact-match figure (`cms_exact_match.png`) is final and raw-to-raw. The other CMS display plots still use the **old JEC-corrected** reference and need a pass:

- [] `make_cms_plots.py` → **compare RAW-to-RAW** — use `FatJet_rawFactor` for the jet pt and `SubJet_rawFactor` + `subJetIdx1/2` for the soft-drop reference (as `exact_match2.py` does).
- [] `cms_recluster.png` — currently vs JEC-corrected `FatJet_pt` (title says "6% mean"); regenerate vs raw pt → should read ≈ 0 .
- [] `cms_softdrop.png` — currently anti-kt-tree grooming vs corrected msoftdrop (−4 GeV); regenerate on the **C/A tree** vs raw subjet mass → should read ≈ 0 .
- [] `cms_lund.png` — regenerate on the C/A tree (currently anti-kt-tree Lund).

Deferred deliberately — the exact-match plot already proves the physics; these are cosmetic consistency fixes for the deck. Tracked in `Projects/Flashjet/Status.md` .

Function reference (added)

```
exclusive_jets_from_history(hist_p1, hist_p2, hist_child, hist_d, mask,
                           n_jets=None, d_cut=None)           # F1
groom_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4, R,
                   z_cut=0.1, beta=0.0, mu=None)             # F2 (soft-drop / mMDT / mass-drop)
lund_coordinates_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4) # F3 -> (B,J,S,6)

# helpers
_resolve_roots(...)    # pointer-jump roots of an arbitrary sub-forest
_jet_roots(...)        # per-jet root pseudojet id, beam-merge order
_pseudojet_p4(...)     # E-scheme 4-momentum of every pseudojet, keyed by id
_dense_number_roots(...) # compact per-particle jet index
```

API surface: `ClusterOutput.exclusive_jets`, `.groomed_jets`, `.mass_drop`,
`.lund_coordinates`, plus a new `.mask` field to map slots → ids.

Reproducing everything

```
# on lxplus, env b_hive (torch 2.5.1)
cd /eos/home-c/cgupta/flashjet/plots/2026-07-13-substructure
micromamba run -n b_hive python make_plots.py           # F1/F2/F3 toy justification + parity
micromamba run -n b_hive python make_paper_plots.py      # anti-kt/SoftDrop/Lund paper figures
micromamba run -n b_hive python make_cms_plots.py        # real CMS UL18 QCD (needs qcd_jmenano_150x.root)

# tests
cd /eos/home-c/cgupta/flashjet/FlastJetDemo
micromamba run -n b_hive python -m pytest -q            # 85 passed, 13 skipped
```

Fixed seed `20260713` throughout. Papers catalogued in `References/Flashjet/papers.md`.