

GPU substructure for flashjet

kt / C-A jet substructure on the merge history

exclusive jets · soft-drop grooming · Lund coordinates

Chirayu Gupta — for Alexandre De Moor

2026-07-13 · branch `benchmarking` (working tree, not committed)

Every claim below is closed against an independent reference and/or the defining paper. All plots regenerable: `plots/2026-07-13-substructure/`.

Where this sits in flashjet

flashjet clusters padded `(B, N, 4)` torch tensors into jets on the GPU, staying on-device for ML training loops. Generalized-kt:

$$d_{ij} = \min(k_{t,i}^{2p}, k_{t,j}^{2p}) \frac{\Delta R_{ij}^2}{R^2}, \quad p = -1 \text{ (anti-} k_t), 0 \text{ (C/A)}, +1 \text{ (} k_t)$$

The key object Alex's kernels already produce: the *merge history*. Every recombination step records four arrays — `hist_p1`, `hist_p2` (the two inputs), `hist_child` (the output id), `hist_d` (the distance at which they merged). That history *is* the full binary clustering tree.

Our contribution = read that tree. No new kernels, no changes to the clustering path. All three features are pure-torch post-reads of `(hist_p1, hist_p2, hist_child, hist_d)`, so they run unchanged on CPU or CUDA and add negligible cost.

What was already there vs. what we added

Pre-existing (Alex)

- `history.py`: `jet_idx_from_history`, `_resolve_parents`, `_decode`, `splitting_scales_from_history`, `Triton`, `_decode_triton`
- `reference.py`: single-event NumPy ground truth (`cluster_event`)
- `nn_reference.py`: nearest-neighbour reference
- `conftest.py`: `random_event` fixtures
- The kernels + the merge-history recording

The `_ref_*` walkers inside `test_substructure.py` are also ours — independent naive NumPy tree-walks used only to pin the fast implementations.

Added (this work) — in `history.py`

- **F1** `exclusive_jets_from_history`
- **F2** `groom_from_history`
- **F3** `lund_coordinates_from_history`
- helpers: `_resolve_roots`, `_jet_roots`, `_pseudojet_p4`, `_dense_number_roots`
- `api.py`: `ClusterOutput`.{`exclusive_jets`, `groomed_jets`, `mass_drop`, `lund_coordinates`} + a `mask` field
- `tests/test_substructure.py` (new)

The three features

	Feature	Function	Implements
F1	Exclusive jets (kt)	<code>exclusive_jets_from_history(...)</code>	kt algorithm — Fastjet manual [1111.6097], [0802.1189]
F2	Grooming (C/A)	<code>groom_from_history(...)</code>	Soft Drop [1402.2657] · mMDT [1307.0007] · mass-drop [0802.2470]
F3	Lund coordinates	<code>lund_coordinates_from_history(...)</code>	Primary Lund plane [1807.04758]

F1 — undo the last merges of the sequence to expose exactly `n_jets` (or a `d_cut`) subjects; reduces to the inclusive jets at the trivial cut.

F2 — walk each jet down the harder branch, dropping the softer prong until $z > z_{\text{cut}}(\Delta R/R)^\beta$ (soft-drop) — the $O(\log_2 n)$ declustering.


F3 — for every primary split emit $(z, \Delta R, k_t, \ln 1/\Delta R, \ln k_t, d)$; the d -channel ties **exactly** to `splitting_scales` .

Design: everything is a cheap post-read


All three reuse two batched **pointer-jumping** primitives that resolve parents / roots of the tree in $O(\log N)$ gathers (`_resolve_parents` , `_resolve_roots`) — no Python-side per-jet loop.

correctness + cost: decoders match an independent walk exactly, and cost a fraction of the clustering

groomed mass parity, 150 events
max $|\text{diff}| = 1.71\text{e-}13$ GeV



features are cheap post-reads of the history



Left: groomed mass from F2 vs an independent per-event NumPy declustering — agree to $\max |\Delta| = 1.71 \times 10^{-13}$ GeV (float64 round-off). **Right:** per-event CPU cost — the decoders are **10-100× cheaper** than the clustering they read; they are not on the critical path.

How the toys were generated

(the "simulation" — none of it pre-existing; flashjet only clusters, it does not generate events)

The toy event generators (written in the plot scripts)

flashjet takes 4-vectors in and returns jets — it has **no event generator**. To exercise the features I wrote small, transparent toys *outside the repo*, in the plot scripts.

`make_plots.py` — **two single-fat-jet samples at $p_T \approx 500$ GeV, $R = 0.8$:**

- **QCD-like:** one hard collinear core (92% of p_T , $\sigma_{y,\phi} = 0.06$) + soft wide-angle radiation (8%, $\sigma = 0.30$). Each "spray" = n massive pions with exponential p_T fractions.
- **W-like:** two hard prongs with pair mass $m = \sqrt{z(1-z)} p_T$, $\Delta R = 80.4$ GeV ($z \in [0.30, 0.45]$), random orientation, + 6% soft contamination.

```
def _spray(rng, n, y0, phi0, pt_total, sigma):    # n pions around (y0,phi0)
    frac = rng.exponential(1., n); frac /= frac.sum()
    pt = pt_total * frac
    y, phi = rng.normal(y0, sigma, n), rng.normal(phi0, sigma, n)
    mt = np.sqrt(M_PI**2 + pt**2)
    return np.stack([pt*np.cos(phi), pt*np.sin(phi), mt*np.sinh(y), mt*np.cosh(y)], 1)
```

1500 events per sample. Purpose: QCD vs W is a *known* truth, so any observable that fails to separate them (or lands off the analytic curve) would be a real bug.

The toy parton shower (for the paper-figure closures)

`make_paper_plots.py` builds a **primary, fixed-coupling, leading-log** shower — emissions sampled **uniformly in the Lund triangle** with density $\bar{\alpha}$ per unit $(\ln 1/\theta, \ln k_t)$ area, off a hard spine at $(y, \phi) = (0, \pi)$:

$$\theta = e^{-u}, \quad k_t = e^v, \quad z = \frac{k_t}{p_{T0} \theta}, \quad \text{keep } v \leq \ln(p_{T0}/2) - u \quad (z \leq \tfrac{1}{2}), \quad k_t > k_t^{\min}$$

This is **exactly the semi-classical picture the papers' analytic predictions are derived in** — which is what makes the closure *quantitative*, not just qualitative.


- **Jet areas:** one hard event + a dense grid of infinitely-soft **ghosts** ($p_T = 10^{-8}$); the ghosts each algorithm sweeps up trace its catchment area (the anti-kt-paper construction).

$\bar{\alpha} = 0.25, p_{T0} = 1 \text{ TeV}, R = 0.4$, 20 000 showers. Everything is re-runnable with a fixed seed (`20260713`).

Correctness

each feature reproduces the signature figure of the paper it implements


anti-kt jet shapes — reproduces [0802.1189] Fig. 1



Same event + soft-ghost grid, $R = 1$. **kt** and **C/A** give ragged, area-fluctuating jets; **anti-kt** gives rigid **circles** around each hard particle — the defining result of the anti-kt paper, here from flashjet's own clustering.

F1 — kt substructure separates 2-prong from QCD [0802.1189, 1111.6097]

kt-based substructure: hard two-prong structure separates from QCD [arXiv:0802.1189, 1111.6097]



Left — $\sqrt{d_{12}}$ (`splitting_scales_from_history`):
 kt scale of the last merge. W-like peaks at $\sim m_W/2$, QCD sits low.

Right — exclusive 2-subject z
 (`exclusive_jets_from_history` , $n_{\text{jets}} = 2$):
 W-like balanced ($z \approx 0.35$), QCD lopsided ($z \rightarrow 0$) — the tagging split.

F2 — soft-drop z_g vs the analytic prediction [1402.2657]

The soft-drop momentum fraction z_g from


`groom_from_history`

($z_{\text{cut}} = 0.1$, $\beta = 0$) tracks the **leading-log QCD prediction**

$$p(z_g) = \frac{1/z_g}{\ln(1/2z_{\text{cut}})}$$

across the **entire range**, with no free parameters. 72% of jets tagged.

This is the **decisive grooming-correctness plot**: the target curve is *the right answer*, so lying on it is unambiguous closure.




F2 — groomed mass ordering in β [1402.2657 Figs. 3-4]

Groomed mass $\rho = m^2/(p_T^2 R^2)$ for $\beta = 0, 1, 2$ vs ungroomed.


- Grooming shifts the spectrum to **lower** ρ (removes soft-wide radiation).
- **Smaller β grooms harder**: the $\beta = 0$ curve reaches furthest left.

Reproduces the β -ordering of the Soft Drop paper.



F3 — primary Lund plane [1807.04758]

F3 lund_coordinates_from_history -- primary Lund plane (C/A, $R=0.8$) [arXiv:1807.04758]




`lund_coordinates_from_history` ($C/A, R = 0.8$). **QCD** fills the soft-collinear region smoothly. **W-like** shows the same background **plus an isolated hard-splitting spot** exactly where the 2-body m_W decay must sit (red ★ = predicted position).

F3 — Lund-triangle closure [1807.04758 Fig. 2]

Emissions sampled **uniformly** in the Lund triangle (input $\bar{\alpha} = 0.25/\text{unit area}$) are recovered by `lund_coordinates_from_history` as a **flat interior plateau** bounded by the three physical edges ($z = \frac{1}{2}$, k_t cutoff, $\theta_{\max}$).

- Interior density **0.17 ± 0.02** — flat to $\sim 12\%$; uniform input recovered.
- The offset below 0.25 is genuine **wide-angle reclustering migration** (C/A reassigns some wide emissions), **not** a bug — it is a property of the observable, reproduced honestly.



F3 runs on **real CMS data** — primary Lund plane of 60 285 QCD jets

Not a toy: **PF candidates** from CMS UL18 QCD **JMENano**

(the format that carries the

PFCand + FatJetPFCand constituent map), grouped per

AK8 jet, clustered by our anti-kt

$R = 0.8$, read by **F3**.

The full physical structure of [1807.04758] emerges

straight from detector-level simulation:

the hard-collinear perturbative ridge, the soft plateau,


and the three kinematic edges — with

no toy input.

The pipeline was also validated jet-by-jet vs CMS AK8 p_T / m_{SD} (omitted here — a

known raw-vs-PUPPI offset, accounted for by a per-jet rescale; the

algorithm is correct).



Validation ladder & test suite

Each feature is pinned at every level, top to bottom:

1. **FastJet** (where applicable) → **NumPy reference** (`cluster_event`) — the single-event ground truth
2. **torch backend** matches the reference
3. **Triton kernels** match the torch backend (`decode=False` parity)
4. **each feature** matches an **independent naive NumPy tree-walk** (`_ref_*` in `test_substructure.py`)
5. **physics anchors**: z_g vs analytic $1/z$; $\sqrt{d_{12}} \rightarrow m_W/2$; Lund hard-split position; groomed-mass parity to 10^{-13} GeV

```
micromamba run -n b_hive python -m pytest -q → 85 passed, 13 skipped (CUDA-only)
```

The 13 skips are the Triton `decode=False` parity tests, which need a GPU node.

Status & next steps

Done

- F1/F2/F3 implemented as pure-torch reads of the merge history — no kernel changes
- Independent-reference parity to float64 round-off; full suite green (85 passed)
- Reproduced the signature figures of anti-kt, Soft Drop, and Lund papers
- F3 demonstrated end-to-end on real CMS UL18 QCD PF candidates

Open / for Alex

- Review the working-tree changes — **not committed** to the repo, per instruction
(`src/flashjet/{history,api,__init__}.py` , `README.md` , `tests/test_substructure.py`)
- GPU-node follow-up: the Triton `decode=False` parity tests still need CUDA
- Optional: expose secondary Lund planes; wire the exclusive-jet API into a tagging demo

Repo: `flashjet/FlastJetDemo` (branch `benchmarking`). Plots + scripts:

`plots/2026-07-13-substructure/` . Papers: `References/Flashjet/papers.md` .

Backup

Function reference (added)

```
exclusive_jets_from_history(hist_p1, hist_p2, hist_child, hist_d, mask,
                           n_jets=None, d_cut=None)           # F1
groom_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4, R,
                  z_cut=0.1, beta=0.0, mu=None)              # F2 (soft-drop / mMDT / mass-drop)
lund_coordinates_from_history(hist_p1, hist_p2, hist_child, hist_d, mask, p4) # F3 -> (B,J,S,6)

# helpers
_resolve_roots(...)      # pointer-jump roots of an arbitrary sub-forest
_jet_roots(...)          # per-jet root pseudojet id, beam-merge order
_pseudojet_p4(...)       # E-scheme 4-momentum of every pseudojet, keyed by id
_dense_number_roots(...) # compact per-particle jet index
```

API surface: `ClusterOutput.exclusive_jets`, `.groomed_jets`, `.mass_drop`,
`.lund_coordinates`, plus a new `.mask` field to map slots → ids.

Reproducing everything

```
# on lxplus, env b_hive (torch 2.5.1)
cd /eos/home-c/cgupta/flashjet/plots/2026-07-13-substructure
micromamba run -n b_hive python make_plots.py          # F1/F2/F3 toy justification + parity
micromamba run -n b_hive python make_paper_plots.py    # anti-kt/SoftDrop/Lund paper figures
micromamba run -n b_hive python make_cms_plots.py      # real CMS UL18 QCD (needs qcd_jmenano_150x.root)

# tests
cd /eos/home-c/cgupta/flashjet/FlastJetDemo
micromamba run -n b_hive python -m pytest -q           # 85 passed, 13 skipped
```

Fixed seed `20260713` throughout. Papers catalogued in `References/Flashjet/papers.md`.